


The background features a stylized profile of a human head facing right. The head is composed of a blue wireframe mesh. Inside the head, there are glowing blue circuit lines and binary code (0s and 1s) floating in the air, suggesting a digital or artificial intelligence theme. The overall color scheme is dark blue with glowing blue highlights.

Machine Learning, Artificial Intelligence, and Big Data Analytics (IL, 4th Semester)

Lecture 4

Recap of previous lecture (1/4)

- *Generalization error* is a common issue to deal with in ML
- It can be decomposed into three terms: *Bias*, *Variance*, and *immutable error*
- High bias → Cannot capture the structure of the data → underfit
- High variance → Too sensitive to the training data → overfit



Recap of previous lecture (2/4)

- Techniques to keep the underfit/overfit under control:
- Split Training-Test



- Cross validation




- Bootstrapping: Resampling with replacement
- Bagging: Bootstrapping multiple times and ensemble of multiple models

Recap of previous lecture (3/4)

- Regression: RMSE, MAE, MAPE, Rsquare
- Classification: Confusion Matrix → accuracy, precision, recall, FP Rate, FN Rate, F1-score, ROC, AUC...

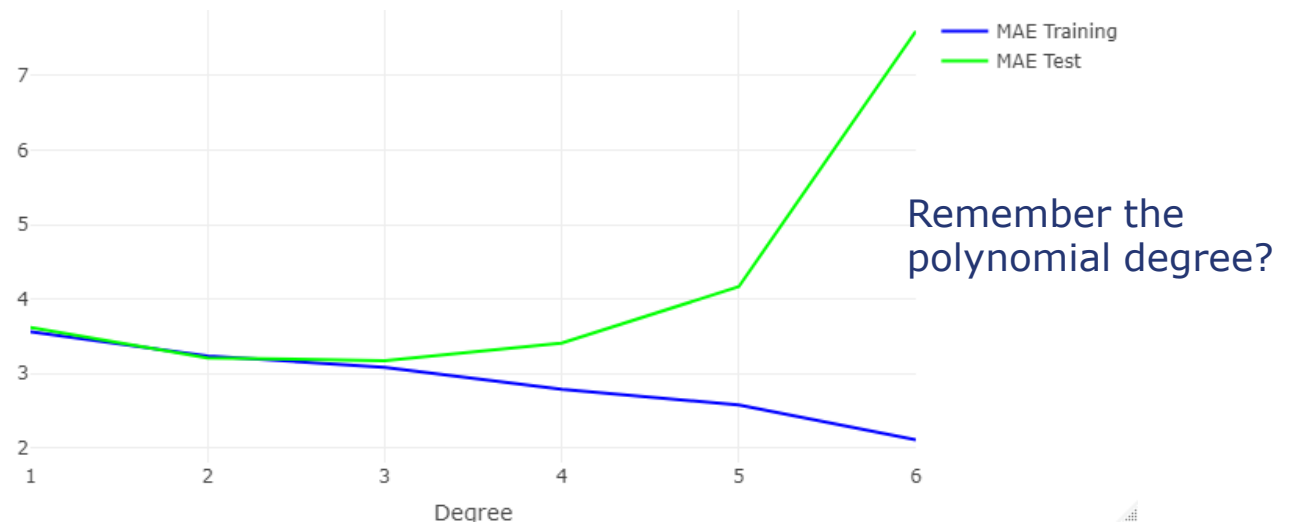
Metric	Formula	Meaning	Visual look	range								
Accuracy	(#TP+#TN)/#Total	What fraction does it get right	<table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table> / <table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table>	TP	FN	FP	TN	TP	FN	FP	TN	0- <u>1</u>
TP	FN											
FP	TN											
TP	FN											
FP	TN											
Precision	#TP/(#TP+#FP)	When it says 1 how often is it right	<table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table> / <table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table>	TP	FN	FP	TN	TP	FN	FP	TN	0- <u>1</u>
TP	FN											
FP	TN											
TP	FN											
FP	TN											
Recall/ Sensitivity	#TP/(#TP+#FN)	What fraction of 1s does it get right (True Positive Rate – TPR)	<table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table> / <table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table>	TP	FN	FP	TN	TP	FN	FP	TN	0- <u>1</u>
TP	FN											
FP	TN											
TP	FN											
FP	TN											
Specificity	#TN/(#TN+#FP)	What fraction of 0s does it get right (True Negative Rate – TNR)	<table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table> / <table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table>	TP	FN	FP	TN	TP	FN	FP	TN	0- <u>1</u>
TP	FN											
FP	TN											
TP	FN											
FP	TN											
FP Rate	#FP/(#FP+#TN)	What fraction of 0s are called 1s	<table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table> / <table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table>	TP	FN	FP	TN	TP	FN	FP	TN	<u>0</u> -1
TP	FN											
FP	TN											
TP	FN											
FP	TN											
FN Rate	#FN/(#TP+#FN)	What fraction of 1s are called 0s	<table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table> / <table><tr><td>TP</td><td>FN</td></tr><tr><td>FP</td><td>TN</td></tr></table>	TP	FN	FP	TN	TP	FN	FP	TN	<u>0</u> -1
TP	FN											
FP	TN											
TP	FN											
FP	TN											
F1-score	$2 * \frac{precision*recall}{precision+recall}$	How “good” are precision and recall		0- <u>1</u>								



Recap of previous lecture (4/4)

What we learned in the coding session

- Stratification helps keeping the class balance in train/test dataset.
- You can make overly complex models and fit perfectly to the training data, but it is almost never a good idea!
- Most of the ML techniques rely on watching the error in the validation/test set and choosing the simplest model with acceptable performance




Agenda

- Decision trees
 - Characterizing a tree
 - Building a tree
 - Pruning a tree
- Coding part
- Exercise

Recall the goal of a classification task

- Classification: Predict category from one or more input variables
- A classification model is trained by learning the relation between input and output from a labeled dataset. It learns how to predict the output variable.



Training phase



Operational phase



Common classification algorithms

- Decision trees
- K-nearest neighbors
- Naïve Bayes
- Support Vector Machine
- Random Forest
- Artificial Neural Network

- And many others...

Decision trees

- Idea: Split data into subsets with one output class




Decision boundaries

"pure": the input space contains sample that have only one output class

Goal: Split the data into regions that are as pure as possible

Decision trees


Structure



- Hierarchical structure with nodes and directed edges
- Root and internal nodes have test conditions
- Leaf nodes have class labels associated with it
- The classification is done by traversing the tree from top until a leaf


Decision trees

Structure



- **Depth of a node:** number of edges from the root to the node
- **Depth of a tree:** Number of edges of the longest path in the tree
- **Size of a tree:** Number of nodes in a tree


Example




Decision trees

Main principle

- Start with all samples in a node
- Partition such that the subsets are as pure as possible



- Repeat the partitioning



- Until a stopping criterion is satisfied

Decision trees


Main principle

- Start with all samples in a node
- Partition such that the subsets are **as pure as possible**



Question 1. How to partition → Splitting criterion

- Repeat the partitioning





- Until a **stopping criterion** is satisfied

Question 2. When to stop → Stopping criterion

Decision trees

Splitting criterion

- How to measure purity?




What's the best splitting here?


Decision trees

Splitting criterion

- How to measure purity?



X




V

Decision trees

How to measure the purity of the subsets

- **Gini impurity index** = $1 - \sum_j p_j^2$
 - All samples one class:
 $1 - 1^2 = 0$
 - Equally distributed in two classes:
 $1 - (0.5^2 + 0.5^2) = 0.5$
- **Entropy** = $-\sum_j (p_j * \log_2 p_j)$
 - All samples one class:
 $-1 * \log_2 1 = 0$
 - Equally distributed in two classes:
 $-0.5 * \log_2 0.5 - 0.5 * \log_2 0.5 = 1$




Decision trees

When to stop growing the tree

Most common criteria

- A percentage of samples have the same class
- The number of samples reaches a minimum
- Change in purity is smaller than a threshold
- Max tree depth is reached




Example of condition	Parameter in <i>rpart</i>
Class > perc_threshold	-
N_samples < min_samples	<i>Minsplit</i>
Purity gain < threshold	<i>CP</i>
Depth == max_depth	<i>Maxdepth</i>

Decision trees

Or grow it all and “prune” it later

Common approach:

- Grow the entire tree
- Plot the results at different “levels” (e.g., different purity gains or different depth)
- Find the level where crossvalidation error is minimal
- “Prune” the tree!



Example of condition	Parameter in <i>rpart</i>
Class > perc_threshold	-
N_samples < min_samples	<i>Minsplit</i>
Purity gain < threshold	<i>CP</i>
Depth == max_depth	<i>Maxdepth</i>


Decision trees

How to split on different types of variables


- Binary



- Categorical



Multi-way split




Binary split

Decision trees


How to split on different types of variables

- Ordinal

- Like categorical but attention to the order



- Numerical



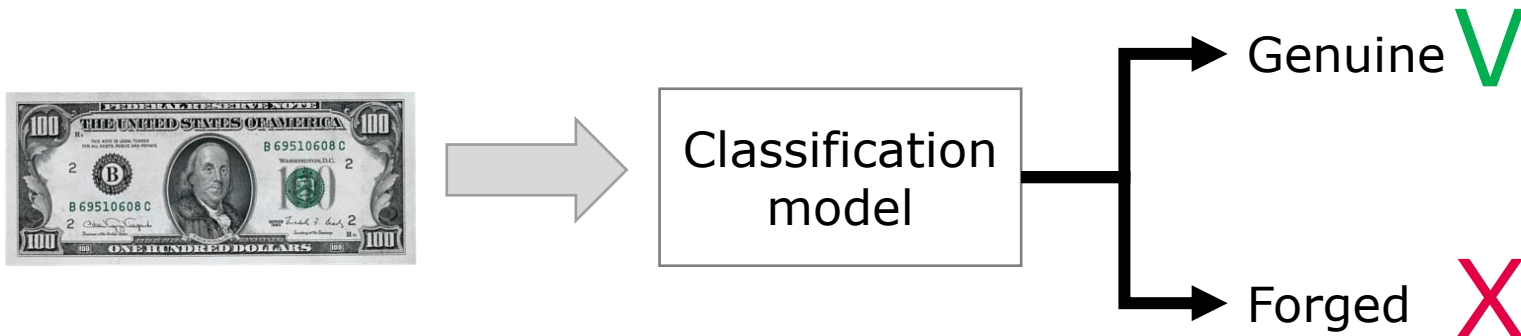
Decision trees

Considerations

- Advantages:
 - Very simple
 - Works with non-linear problems
 - Interpretable
 - Computationally cheap
- There are some drawbacks
 - Unstable (A change in the data can lead to completely different model)
 - Mostly based on heuristic with no solid statistical background
 - Prone to overfit
- Tree ensemble techniques (Random Forest, XGBoost) will be covered in coming lectures.


Decision tree

Practical Example – Detection of forged banknotes





Decision tree

Practical Example – Detection of forged banknotes




Decision tree

Practical Example – Detection of forged banknotes



Training phase



Operational phase

Decision tree


Practical Example – Detection of forged banknotes




Wavelet
Transform



Variance of the wavelet
Skewness of the wavelet
Curtosis of the wavelet
Entropy of the wavelet



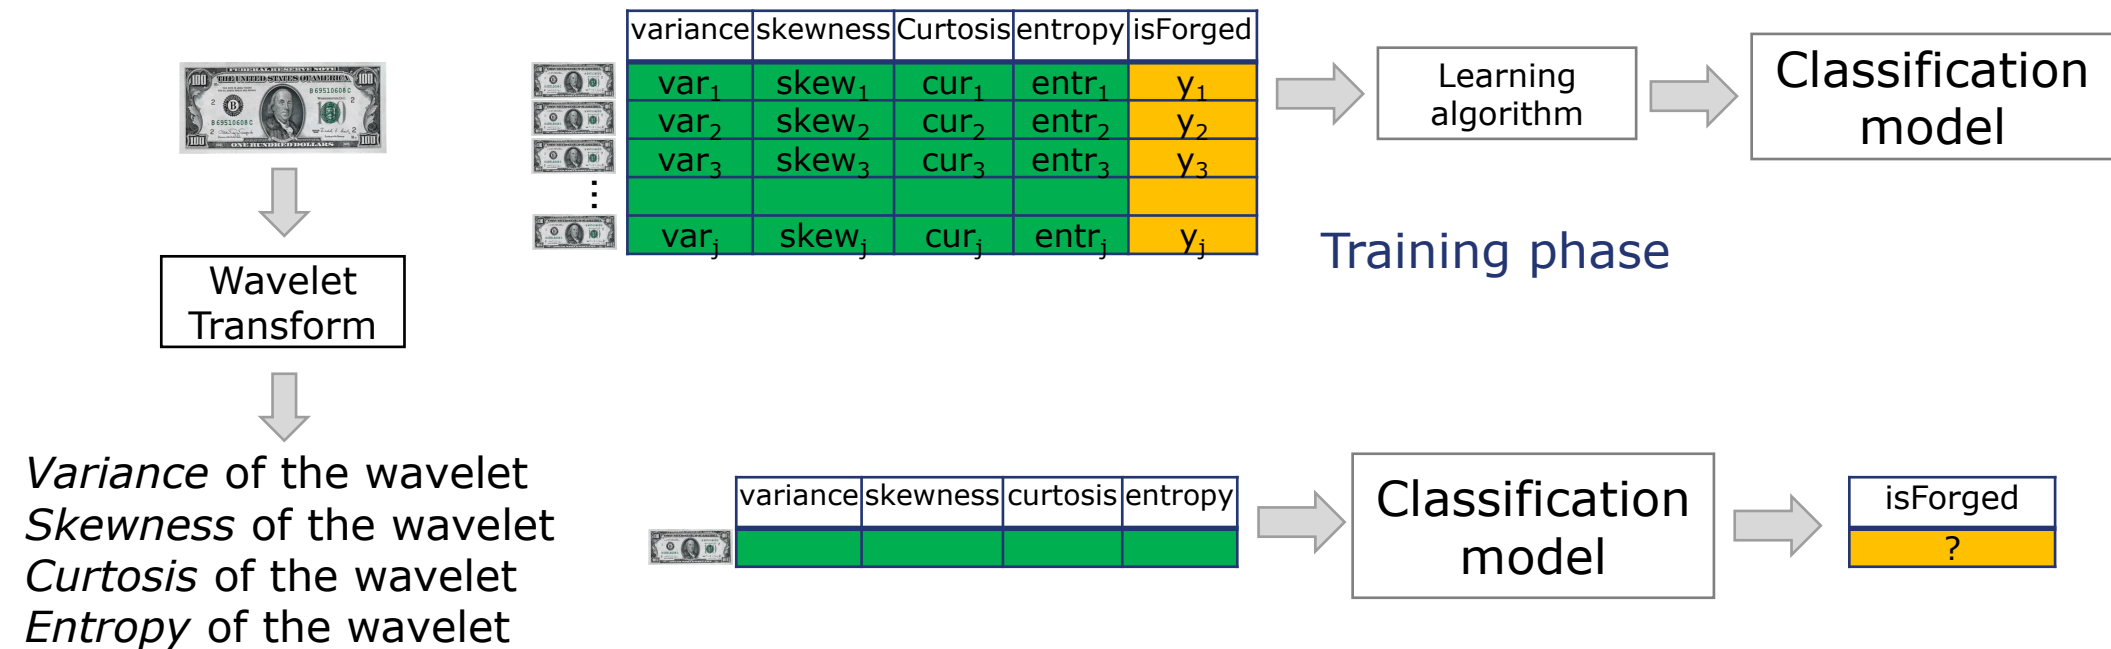
Training phase



Operational phase

Decision tree

Practical Example – Detection of forged banknotes




Let's do it!

Exercise 2 – Due on 27.03.2023 23.59 CET

(late submission +1 week, max 6 points)

- Import the EEG data from the assignment
 - 16 EEG numerical features
 - 1 label (1 = eyes open, 2 = eyes closed)
 - Train and test are already split (*eeg_training.csv*, *eeg_test.csv*)
- Create a classifier that can detect if the patient has open or closed eyelid
- Part 1: Decision Trees
- Part 2: Random Forest, AdaBoost, and XGBoost
- **Attention: This is not the same data that can be found online. Do not copy!**
- Hints: Are there NAs? Is it balanced? Are there useless features? Are there outliers?




Lecture E-LEARNING


- Register to datacamp.com using your IMC email address following [this](#) invite.
- You'll find one assignment already: "Support Vector Machines with R"
- **Due date: March 31st**

Lecture E-LEARNING


Support Vector Machines



Linear



Radial



Polynomial