Machine Learning, Artificial Intelligence, and Big Data Analytics (IL, 4th Semester)

Lecture 8

Where we are (Supervised learning)

- Regression and Classification, generalization error, data splitting, Bootstrapping, Bagging/Boosting, model validation, …

- K-Nearest Neighbors

- Support Vector Machines

- Decision Trees

- Random Forest

- Adaboost

- XGBoost

Coming next (Unsupervised learning)

- Introduction to unsupervised learning
- **Clustering** objective
- Partitioning methods
    - K-means
    - Partitioning Around Medoids (PAM)
- Hierarchical methods
    - Agglomerative (AGNES)
    - Divisive (DIANA)
- Density-based methods
    - DBscan and OPTICS

Agenda

- Introduction to unsupervised learning
- Clustering
- Types of Clustering
- K-means algorithm
- K-means variants
- Validating clustering

Supervised learning

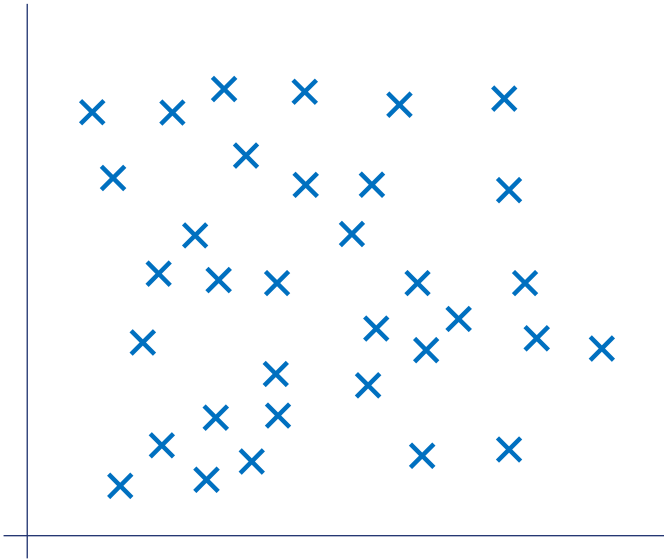Training data is labeled (e.g., green square vs. red circle)
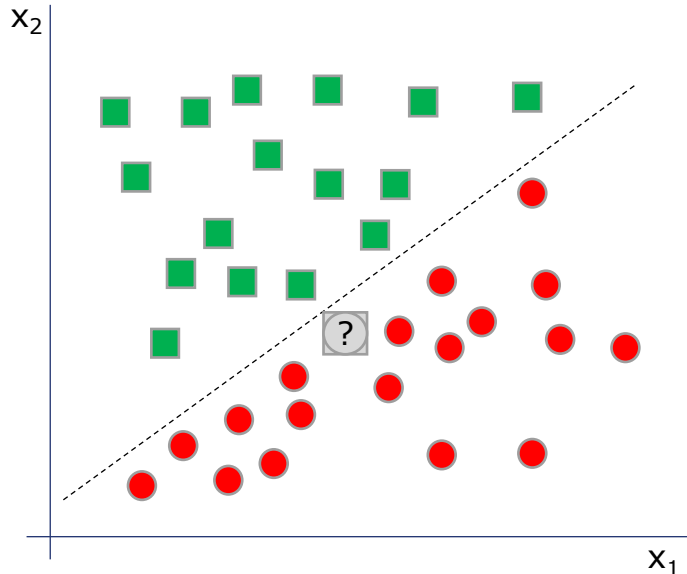
Training set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \cdots, (x^{(m)}, y^{(m)})\}$

GOAL: learn $f(x) \rightarrow y$

Unsupervised learning

Training data is not labeled

Training set: $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots, x^{(m)}\}$

GOAL: find interesting things in data

Examples of unsupervised learning techniques

- Clustering
- Dimensionality reduction
- Anomaly detection
- Association rules mining
- Pattern recognition

Clustering

- Cluster analysis: Given a set of data objects find the proper grouping such that
  - Points in the same groups are similar to each other
  - Points in one group differ from points in other groups
- In other words: Finding <mark>natural groupings</mark> among objects in a dataset
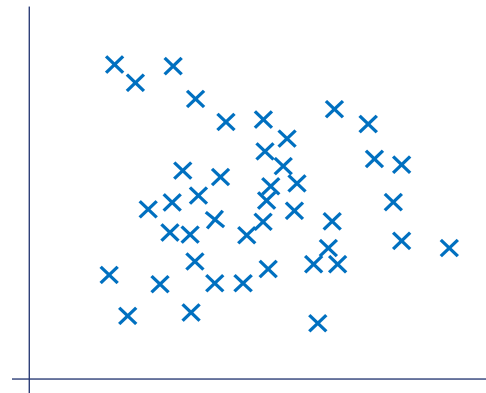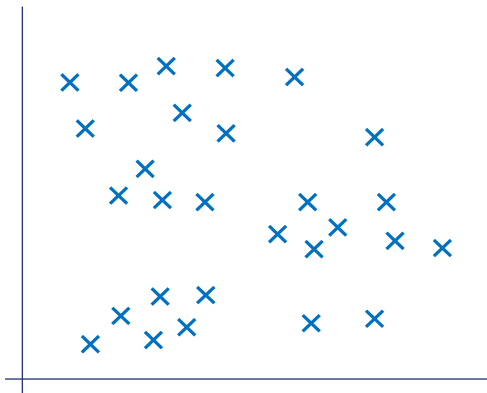
Clustering

- Cluster analysis: Given a set of data objects find the proper grouping such that
  - Points in the same groups are similar to each other
  - Points in one group differ from points in other groups

- In other words: Finding natural groupings among objects in a dataset

Clustering

- Optimal clusters should
  - Maximize similarity within clusters (intra-cluster): *cohesive* within clusters
  - Minimize similarity between clusters (inter-cluster): *distinctive* between cluster

Maximize these distances

Minimize these distances
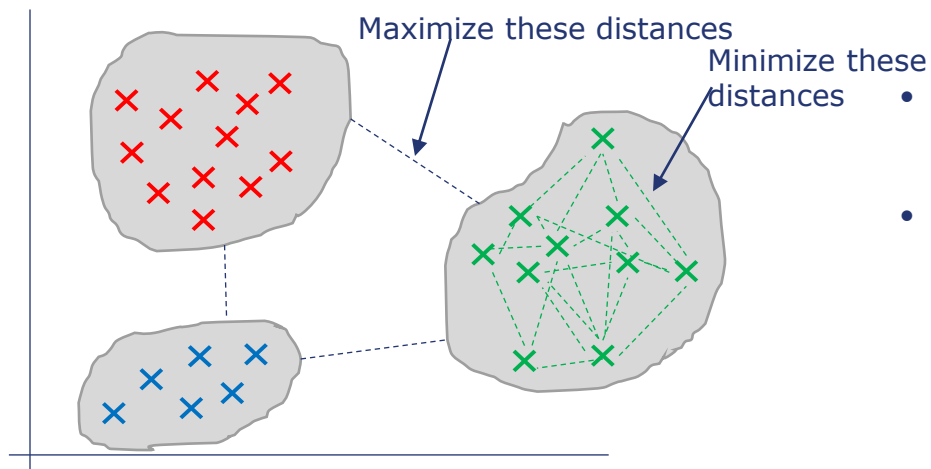
- dissimilarity is expressed in terms of a distance function, typically metric: $d(i, j)$

- The definitions of distance functions are usually rather different for interval-scaled, boolean, categorical, ordinal ratio, and vector variables

Example of clustering

- Biology: Taxonomy of living things: kingdom, phylum, class, order, family, genus and species

- Behavioral: Understanding behavior of the masses

- Information retrieval: Document clustering

- Marketing: Discover distinct groups in their customer bases (customer segments), and use this knowledge to develop targeted marketing programs

- Climate: Understanding earth climate, find patterns

- Mobility: Understanding mobility patterns

- …

The clustering process

Formulate the problem

• Given a set of <…>, I want to find <…>, based on <…>

Select a distance/(dis)similarity measure

• Euclidean, Manhattan, Cosine, Jaccard, …

Select a clustering technique

• K-means, PAM, Diana, DBScan, …

Decide on the number of clusters

Interpret the clusters

• Using Silhouette, Elbow, Dendograms, …

Assess the clustering
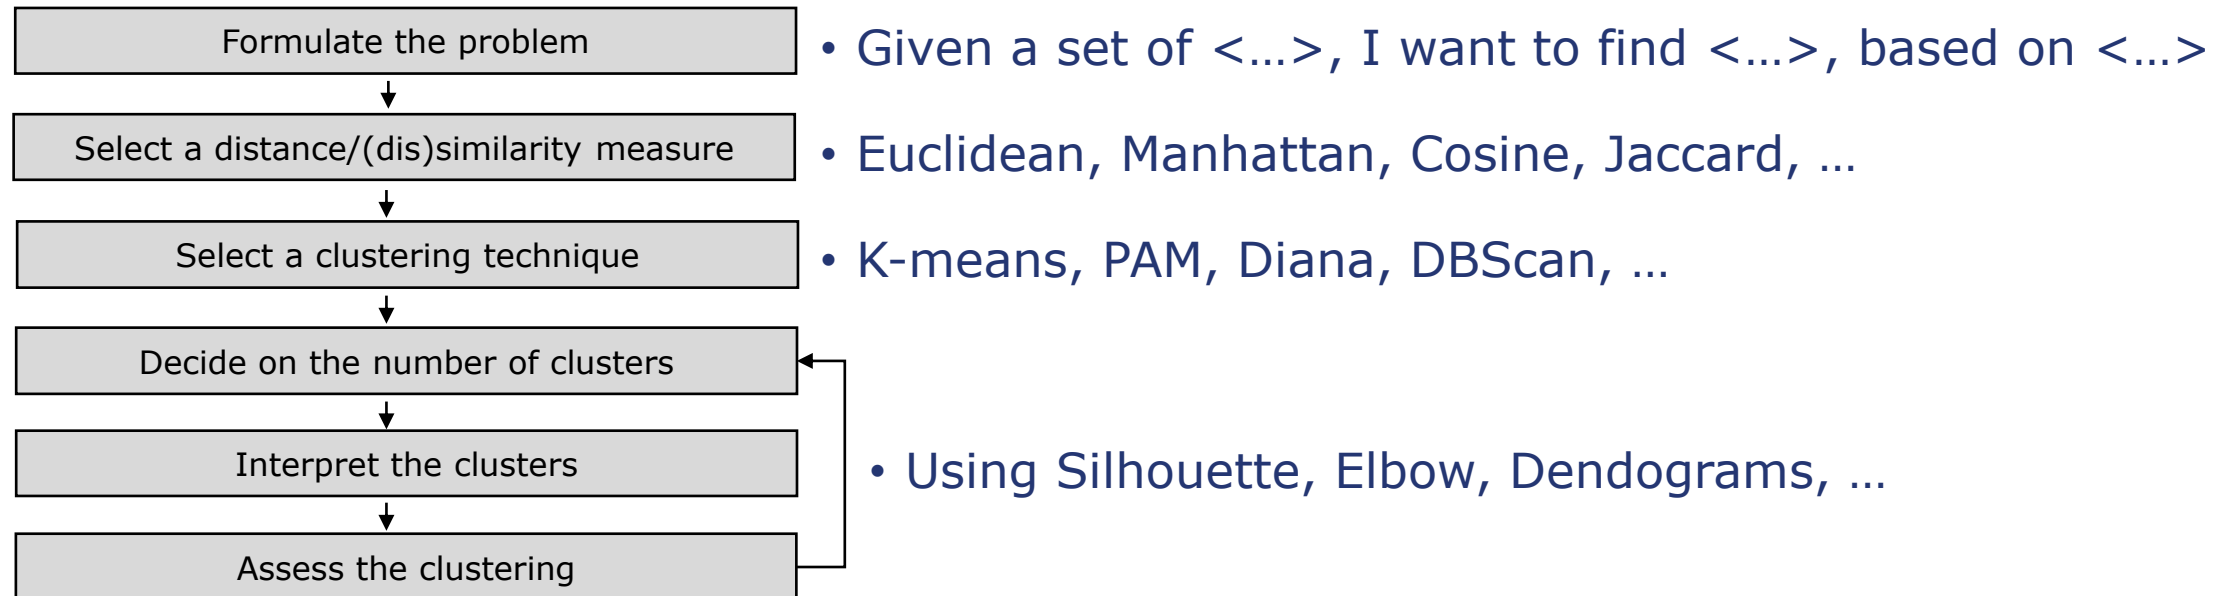
Types of clustering techniques

Types of algorithms

- **Partitional algorithms:**
  - Construct various partitions and then evaluate them by some criterion.
  - Typical methods: k-means, k-medoids/PAM.

- **Hierarchical algorithms:**
  - Create a hierarchical decomposition of the set of objects using some criterion.
  - Typical methods: Diana, Agnes.

- **Density-based algorithms:**
  - Based on Connectivity and density functions
  - Typical methods: DBscan, OPTICS.

- …

Types of clustering techniques

Hard vs. Soft Clustering

- **Hard clustering**
  - Each sample belongs to exactly one cluster
  - For example: An *animal* belong to a *species*

- **Soft clustering**
  - A sample can belong to more than one cluster (probabilistic)
  - For example: *Glasses* belong to *medical aid* and to *fashion item*

Partitioning algorithms

- <u>Partitioning method:</u> Partitioning a database **D** of **n** objects into a set of **k** clusters, such that the sum of intra-cluster squared distances is minimized

- Given *k*, find a partition of *k clusters* that optimizes the chosen partitioning criterion
  - Global optimal: exhaustively enumerate all partitions
  - Heuristic methods: *k-means* and *k-medoids* algorithms
    - <u>*k-means*</u>: Each cluster is represented by the center of the cluster
    - <u>*k-medoids*</u> or PAM (Partition around medoids): Each cluster is represented by one of the objects in the cluster

K-Means algorithm

- Input:
  - K (number of clusters)
  - Training set $\{x^{(1)}, x^{(2)}, x^{(3)}, \cdots, x^{(m)}\}$

- Algorithm
  - Randomly initialize K cluster centroids $\mu_1, \mu_2, \cdots, \mu_K \in \mathbb{R}^n$
  - Repeat{

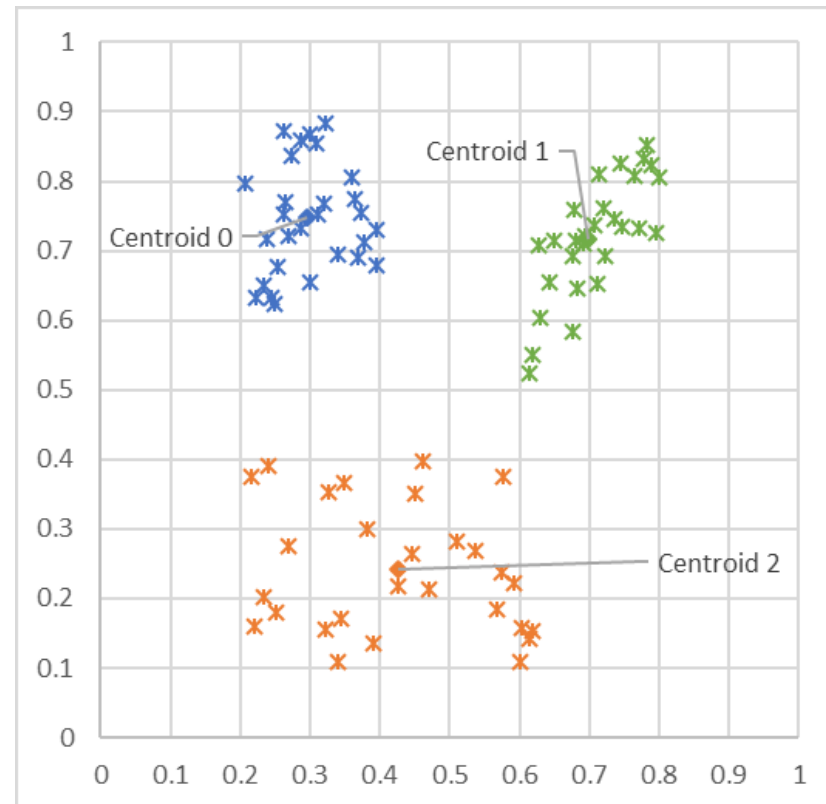|	
for $i = 1$ to m $\qquad c^{(i)} :=$ index (from 1 to K) of cluster centroid closest to $x^{(i)}$	Cluster assignment step
for $k = 1$ to K $\qquad \mu_k :=$ average (mean) of points assigned to cluster k	Centroid update step

  }

K-Means visual example

- Unlabeled Data

- Initialize centroids to random data point locations

- Each point assigned to nearest centroid

- Update centroid locations to avg of assigned points

- Each point assigned to nearest centroid

- Update centroid locations to avg of assigned points

- Each point assigned to nearest centroid

- Converge!

K-Means - How to choose K

- Pick a K based on your understanding of the domain
- Run K-Means
- Examine samples from each cluster
- Adapt K based on what you find
  - If single clusters contain different entities, increase K
  - If entities spread across clusters, decrease K
- OR
- Try multiple values of K and pick the K that maximizes a specified metric

K-Means considerations

- The objective of k-means is to minimize the total sum of the squared distance of every point to its corresponding cluster centroid.

- Finding the global optimum is NP-hard.

- The k-means algorithm converges to a **local** optimum.

- Results can vary based on random seed selection.

- Some seeds can result in poor convergence rate, or convergence to sub-optimal clusterings.

K-Means weaknesses

- Sensitive to outliers
- Needs an initial guess on the number of clusters
- Results depend on the initial seeds
- Can be applied only to objects in a continuous n-dimensional space

Variants of K-means

- Many variants, usually differ by
  - Selection of the initial K
  - Initialization
  - Dissimilarity calculation

- Popular variants
  - K-medians:
    - Cluster center is the median
  - K-medoid:
    - Cluster center is an actual datapoint (not same as k-medians)

- Algorithms are similar to k-means

Evaluating the quality of clustering

- Two type of measures: intrinsic and extrinsic

- Intrinsic: unsupervised, i.e., the ground truth is unavailable
  - Evaluate the goodness of a clustering by considering how well the clusters are separated, and how compact the clusters are (e.g., intracluster similarity, intercluster dissimilarity, the Silhouette coefficient)

- Extrinsic: supervised, i.e., the ground truth is available for at least a subset of data
  - Compare a clustering against the ground truth using certain clustering quality measure

Evaluating the quality of clustering

The "Elbow method"

- SSE (or WSS): Sum of Squared Distances between data points and their assigned cluster's centroid
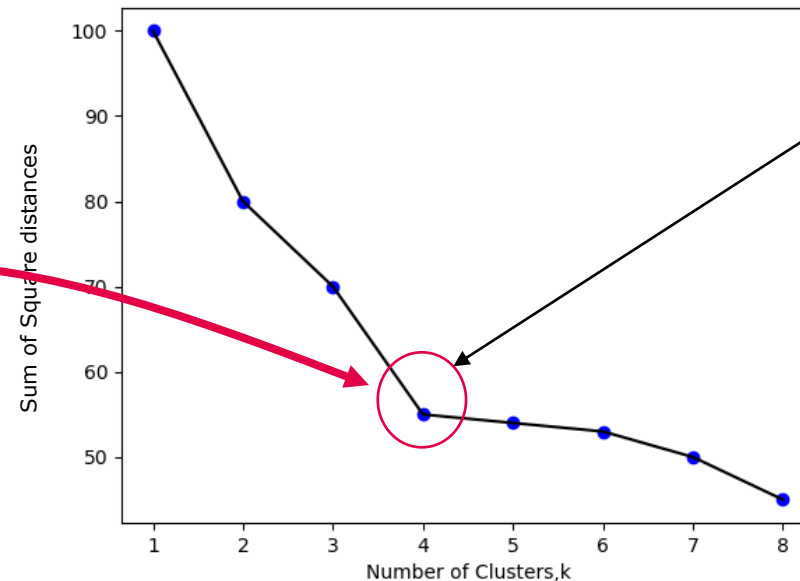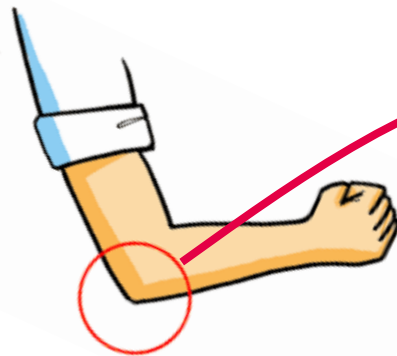
Elbow → n=4

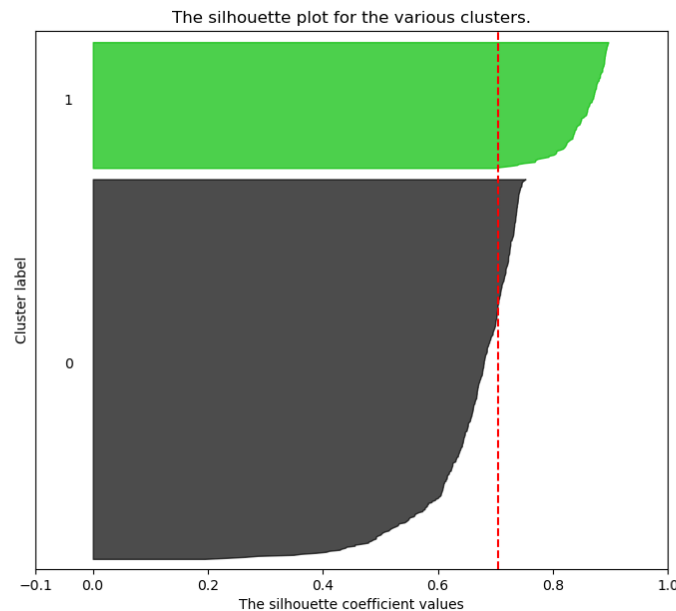# Evaluating the quality of clustering

The silhouette method

- Silhouette of one observation: $s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$

- Where
  - *a(i) = average distance between i and all points in the <u>same</u> cluster*
  - *b(i) = average distance between i and all points in the <u>closest</u> different cluster*
- Interpretation:
  - *s(i) = 1* → observation fits well in its cluster and is far from other clusters
  - *s(i) = 0* → observation is as close to its cluster as to the neighbor cluster
  - *s(i) = -1* → observation is closer to the neighbor cluster than to its own
- Silhouette Score
  - SS=*mean{s(i)}* is a generic measure of cluster cohesiveness and distinction.

Evaluating the quality of clustering

The silhouette method (example)

Silhouette analysis for KMeans clustering on sample data with n_clusters = 2

From: https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html

Evaluating the quality of clustering

The silhouette method (example)

Silhouette analysis for KMeans clustering on sample data with n_clusters = 3

if the values are negatieve that means that the points are closer to another cluster thatn to caster they are asigned ot

From: https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html

Evaluating the quality of clustering

The silhouette method (example)

Silhouette analysis for KMeans clustering on sample data with n_clusters = 4

From: https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html

Evaluating the quality of clustering

The silhouette method (example)

Silhouette analysis for KMeans clustering on sample data with n_clusters = 5

From: https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html

Evaluating the quality of clustering

The silhouette method (example)

Silhouette analysis for KMeans clustering on sample data with n_clusters = 6

From: https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html
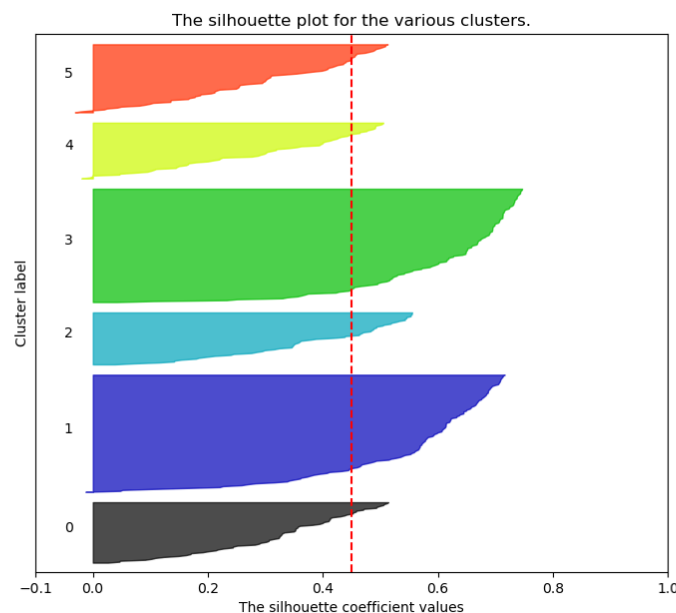
DEADLINE: **May 8th, 09.00** (NO LATE SUBMISSIONS)
(MEGABONUS: **+5**) Do the assignment both in R and Python

- Create a report (**Rmarkdown and/or Jupyter notebook**) addressing the following assignments. The report must contain both code (cells) and results (no need to re-run)!

  image manipulation, ppillow, opencv,

3.1 Mobile phone picture

- Take a picture with your smartphone. The picture must contain a piece of paper with your name on it and some type of background (walls, floor, window, etc.).
- Resize it to a manageable size (e.g., 256x256) either with R or Py
- The goal is to reduce the number of RGB colors by using k-means as in the lecture.
- Pick the k suggested by the elbow mechanism. Try also other k values.
- The report must contain the original pic, the WSS plot (elbow), & the final pictures

example

3.2 Drilling machine

- *drilling.csv* contains 400 operational measurements from a drilling machine.
- The machine can operate in different unknown states.

  - Identify the number of states by using the known clustering techniques
    - K-means (iterate over k → elbow → final clustering)
    - Hierarchical clustering (iterate over linkages → AC/dendrograms → final clustering)
    - Dbscan (kNNdistplot → Eps → final clustering)
    - Optics
  - The report must contain all plots and a final comparison of the different clustering outcomes

- Coding session